

Beginner Level Challenge:

Project: Deploy a Simple Web Server on AWS with Terraform

Objective:

- Use the AWS provider to deploy an EC2 instance.
- Store Terraform state locally.

Steps:

1. Set up a basic AWS provider.
2. Define variables for AMI, instance type, and region.
3. Deploy a t2.micro EC2 instance.
4. Output the instance's public IP address.
5. Use the local backend for storing the state file.

Key Learning Outcomes:

- Understanding providers, resources, variables, and outputs.
- Practice with local state management.

Improvement Suggestions:

- Incorporate variable types and defaults to enhance the flexibility of the configuration.

Intermediate Level Challenge:

Project: Manage Multiple Environments Using Terraform Workspaces

Objective:

- Implement workspaces to manage different environments (development and production).
- Use remote state storage with S3 and enable state locking with DynamoDB.

Steps:

1. Set up AWS CLI for managing the infrastructure.
2. Create S3 and DynamoDB resources for remote state management.
3. Create a Terraform configuration to deploy EC2 instances in different environments using workspaces.
4. Output instance IDs and public IPs for each environment.
5. Use IAM policies to secure the S3 bucket and DynamoDB table.

Key Learning Outcomes:

- Understanding Terraform state management with remote backends.
- Introduction to workspaces for multi-environment support.
- Practice with IAM roles and access controls.

Improvement Suggestions:

- Add lifecycle management configurations (`create_before_destroy`) to handle resource updates and avoid downtime.

Advanced Level Challenge:**Project:**

Automate Infrastructure Scaling with Auto Scaling Group, Load Balancer, and S3 Bucket

Objective:

- Design a scalable architecture with multiple EC2 instances, an Elastic Load Balancer, Auto Scaling Group, and store assets in an S3 bucket.
- Securely manage the infrastructure with Terraform's remote state, workspaces, and state locking.

Steps:

1. Set up AWS CLI and configure S3 and DynamoDB for remote state and state locking.
2. Use modules to create an Auto Scaling Group, Elastic Load Balancer, and an S3 bucket for asset storage.
3. Implement variables and workspaces to switch between different environments (dev, staging, prod).
4. Use version control (Git) to track Terraform changes.
5. Set up an EC2 instance profile with permissions to interact with S3.
6. Ensure proper monitoring and scaling policies in place.

Key Learning Outcomes:

- Handling complex, multi-resource infrastructure with Terraform.
- Advanced state management practices with real-world scaling and monitoring needs.
- Workspaces, version control, and disaster recovery strategies.

Improvement Suggestions:

- Implement advanced Terraform techniques like data sources and conditional expressions to optimize resources.

- Integrate monitoring using CloudWatch or Prometheus, and alarms for Auto Scaling.

General Suggestions for Improvement:

1. Incorporate Best Practices : Emphasize the importance of version control, especially for intermediate and advanced learners.

2. Security Best Practices: For all levels, stress the need for securing state files (especially remote ones) using encryption in S3 and proper IAM roles.

3. Real-World Scenario-Based Labs : Introduce common industry challenges (like disaster recovery, compliance, and team collaboration) into each level's project.

These challenges will provide a progressive learning curve for beginners to experts and can help learners apply Terraform concepts in real-world scenarios.